

11th November 2021: PostgreSQL 14.1, 13.5, 12.9, 11.14, 10.19, and 9.6.24 Released!

[Documentation](#) → [PostgreSQL 9.5](#)Supported Versions: [Current \(14\)](#) / [13](#) / [12](#) / [11](#) / [10](#)Development Versions: [devel](#)Unsupported versions: [9.6](#) / [9.5](#) / [9.4](#) / [9.3](#) / [9.2](#) / [9.1](#) / [9.0](#) / [8.4](#) / [8.3](#) / [8.2](#) / [8.1](#) / [8.0](#) / [7.4](#) / [7.3](#) / [7.2](#) / [7.1](#)

Search the documentation for...



This documentation is for an unsupported version of PostgreSQL.

You may want to view the same page for the [current](#) version, or one of the other supported versions listed above instead.[Prev](#)[Up](#)[PostgreSQL 9.5.25 Documentation](#)

Chapter 9. Functions and Operators

[Next](#)

9.20. Aggregate Functions

Aggregate functions compute a single result from a set of input values. The built-in normal aggregate functions are listed in [Table 9-49](#) and [Table 9-50](#). The built-in ordered-set aggregate functions are listed in [Table 9-51](#) and [Table 9-52](#). Grouping operations, which are closely related to aggregate functions, are listed in [Table 9-53](#). The special syntax considerations for aggregate functions are explained in [Section 4.2.7](#). Consult [Section 2.7](#) for additional introductory information.

Table 9-49. General-Purpose Aggregate Functions

Function	Argument Type(s)	Return Type	Description
<code>array_agg(<i>expression</i>)</code>	any non-array type	array of the argument type	input values, including nulls, concatenated into an array
<code>array_agg(<i>expression</i>)</code>	any array type	same as argument data type	input arrays concatenated into array of one higher dimension (inputs must all have same dimensionality, and cannot be empty or null)
<code>avg(<i>expression</i>)</code>	smallint, int, bigint, real, double precision, numeric, or interval	numeric for any integer-type argument, double precision for a floating-point argument, otherwise the same as the argument data type	the average (arithmetic mean) of all non-null input values
<code>bit_and(<i>expression</i>)</code>	smallint, int, bigint, or bit	same as argument data type	the bitwise AND of all non-null input values, or null if none
<code>bit_or(<i>expression</i>)</code>	smallint, int, bigint, or bit	same as argument data type	the bitwise OR of all non-null input values, or null if none
<code>bool_and(<i>expression</i>)</code>	bool	bool	true if all input values are true, otherwise false
<code>bool_or(<i>expression</i>)</code>	bool	bool	true if at least one input value is true, otherwise false
<code>count(*)</code>		bigint	number of input rows
<code>count(<i>expression</i>)</code>	any	bigint	number of input rows for which the value of <i>expression</i> is not null
<code>every(<i>expression</i>)</code>	bool	bool	equivalent to <code>bool_and</code>
<code>json_agg(<i>expression</i>)</code>	any	json	aggregates values, including nulls, as a JSON array
<code>jsonb_agg(<i>expression</i>)</code>	any	jsonb	aggregates values, including nulls, as a JSON array
<code>json_object_agg(<i>name</i>, <i>value</i>)</code>	(any, any)	json	aggregates name/value pairs as a JSON object; values can be null, but not names
<code>jsonb_object_agg(<i>name</i>, <i>value</i>)</code>	(any, any)	jsonb	aggregates name/value pairs as a JSON object; values can be null, but not names
<code>max(<i>expression</i>)</code>	any numeric, string, date/time, network, or enum type, or arrays of these types	same as argument type	maximum value of <i>expression</i> across all non-null input values
<code>min(<i>expression</i>)</code>	any numeric, string, date/time, network, or enum type, or arrays of these types	same as argument type	minimum value of <i>expression</i> across all non-null input values
<code>string_agg(<i>expression</i>, <i>delimiter</i>)</code>	(text, text) or (bytea, bytea)	same as argument types	non-null input values concatenated into a string, separated by delimiter

Function	Argument Type(s)	Return Type	Description
<code>sum(<i>expression</i>)</code>	<code>smallint</code> , <code>int</code> , <code>bigint</code> , <code>real</code> , <code>double precision</code> , <code>numeric</code> , <code>interval</code> , or <code>money</code>	<code>bigint</code> for <code>smallint</code> or <code>int</code> arguments, <code>numeric</code> for <code>bigint</code> arguments, otherwise the same as the argument data type	sum of <i>expression</i> across all non-null input values
<code>xmlagg(<i>expression</i>)</code>	<code>xml</code>	<code>xml</code>	concatenation of non-null XML values (see also Section 9.14.1.7)

It should be noted that except for `count`, these functions return a null value when no rows are selected. In particular, `sum` of no rows returns null, not zero as one might expect, and `array_agg` returns null rather than an empty array when there are no input rows. The `coalesce` function can be used to substitute zero or an empty array for null when necessary.

Note: Boolean aggregates `bool_and` and `bool_or` correspond to standard SQL aggregates `every` and `any` or `some`. As for `any` and `some`, it seems that there is an ambiguity built into the standard syntax:

```
SELECT b1 = ANY((SELECT b2 FROM t2 ...)) FROM t1 ...;
```

Here `ANY` can be considered either as introducing a subquery, or as being an aggregate function, if the subquery returns one row with a Boolean value. Thus the standard name cannot be given to these aggregates.

Note: Users accustomed to working with other SQL database management systems might be disappointed by the performance of the `count` aggregate when it is applied to the entire table. A query like:

```
SELECT count(*) FROM sometable;
```

will require effort proportional to the size of the table: PostgreSQL will need to scan either the entire table or the entirety of an index which includes all rows in the table.

The aggregate functions `array_agg`, `json_agg`, `jsonb_agg`, `json_object_agg`, `jsonb_object_agg`, `string_agg`, and `xmlagg`, as well as similar user-defined aggregate functions, produce meaningfully different result values depending on the order of the input values. This ordering is unspecified by default, but can be controlled by writing an `ORDER BY` clause within the aggregate call, as shown in [Section 4.2.7](#). Alternatively, supplying the input values from a sorted subquery will usually work. For example:

```
SELECT xmlagg(x) FROM (SELECT x FROM test ORDER BY y DESC) AS tab;
```

But this syntax is not allowed in the SQL standard, and is not portable to other database systems.

[Table 9-50](#) shows aggregate functions typically used in statistical analysis. (These are separated out merely to avoid cluttering the listing of more-commonly-used aggregates.) Where the description mentions *N*, it means the number of input rows for which all the input expressions are non-null. In all cases, null is returned if the computation is meaningless, for example when *N* is zero.

Table 9-50. Aggregate Functions for Statistics

Function	Argument Type	Return Type	Description
<code>corr(<i>Y</i>, <i>X</i>)</code>	<code>double precision</code>	<code>double precision</code>	correlation coefficient
<code>covar_pop(<i>Y</i>, <i>X</i>)</code>	<code>double precision</code>	<code>double precision</code>	population covariance
<code>covar_samp(<i>Y</i>, <i>X</i>)</code>	<code>double precision</code>	<code>double precision</code>	sample covariance
<code>regr_avgx(<i>Y</i>, <i>X</i>)</code>	<code>double precision</code>	<code>double precision</code>	average of the independent variable ($\text{sum}(X) / N$)
<code>regr_avgy(<i>Y</i>, <i>X</i>)</code>	<code>double precision</code>	<code>double precision</code>	average of the dependent variable ($\text{sum}(Y) / N$)
<code>regr_count(<i>Y</i>, <i>X</i>)</code>	<code>double precision</code>	<code>bigint</code>	number of input rows in which both expressions are nonnull
<code>regr_intercept(<i>Y</i>, <i>X</i>)</code>	<code>double precision</code>	<code>double precision</code>	y-intercept of the least-squares-fit linear equation determined by the (<i>x</i> , <i>y</i>) pairs
<code>regr_r2(<i>Y</i>, <i>X</i>)</code>	<code>double precision</code>	<code>double precision</code>	square of the correlation coefficient
<code>regr_slope(<i>Y</i>, <i>X</i>)</code>	<code>double precision</code>	<code>double precision</code>	slope of the least-squares-fit linear equation determined by the (<i>x</i> , <i>y</i>) pairs

Function	Argument Type	Return Type	Description
<code>regr_sxx(<i>Y</i>, <i>X</i>)</code>	double precision	double precision	$\text{sum}(X^2) - \text{sum}(X)^2 / N$ ("sum of squares" of the independent variable)
<code>regr_sxy(<i>Y</i>, <i>X</i>)</code>	double precision	double precision	$\text{sum}(X * Y) - \text{sum}(X) * \text{sum}(Y) / N$ ("sum of products" of independent times dependent variable)
<code>regr_syy(<i>Y</i>, <i>X</i>)</code>	double precision	double precision	$\text{sum}(Y^2) - \text{sum}(Y)^2 / N$ ("sum of squares" of the dependent variable)
<code>stddev(<i>expression</i>)</code>	smallint, int, bigint, real, double precision, or numeric	double precision for floating-point arguments, otherwise numeric	historical alias for <code>stddev_samp</code>
<code>stddev_pop(<i>expression</i>)</code>	smallint, int, bigint, real, double precision, or numeric	double precision for floating-point arguments, otherwise numeric	population standard deviation of the input values
<code>stddev_samp(<i>expression</i>)</code>	smallint, int, bigint, real, double precision, or numeric	double precision for floating-point arguments, otherwise numeric	sample standard deviation of the input values
<code>variance(<i>expression</i>)</code>	smallint, int, bigint, real, double precision, or numeric	double precision for floating-point arguments, otherwise numeric	historical alias for <code>var_samp</code>
<code>var_pop(<i>expression</i>)</code>	smallint, int, bigint, real, double precision, or numeric	double precision for floating-point arguments, otherwise numeric	population variance of the input values (square of the population standard deviation)
<code>var_samp(<i>expression</i>)</code>	smallint, int, bigint, real, double precision, or numeric	double precision for floating-point arguments, otherwise numeric	sample variance of the input values (square of the sample standard deviation)

Table 9-51 shows some aggregate functions that use the *ordered-set aggregate* syntax. These functions are sometimes referred to as "inverse distribution" functions.

Table 9-51. Ordered-Set Aggregate Functions

Function	Direct Argument Type(s)	Aggregated Argument Type(s)	Return Type	Description
<code>mode() WITHIN GROUP (ORDER BY <i>sort_expression</i>)</code>		any sortable type	same as sort expression	returns the most frequent input value (arbitrarily choosing the first one if there are multiple equally-frequent results)
<code>percentile_cont(<i>fraction</i>) WITHIN GROUP (ORDER BY <i>sort_expression</i>)</code>	double precision	double precision or interval	same as sort expression	continuous percentile: returns a value corresponding to the specified fraction in the ordering, interpolating between adjacent input items if needed
<code>percentile_cont(<i>fractions</i>) WITHIN GROUP (ORDER BY <i>sort_expression</i>)</code>	double precision[]	double precision or interval	array of sort expression's type	multiple continuous percentile: returns an array of results matching the shape of the <i>fractions</i> parameter, with each non-null element replaced by the value corresponding to that percentile
<code>percentile_disc(<i>fraction</i>) WITHIN GROUP (ORDER BY <i>sort_expression</i>)</code>	double precision	any sortable type	same as sort expression	discrete percentile: returns the first input value whose position in the ordering equals or exceeds the specified fraction
<code>percentile_disc(<i>fractions</i>) WITHIN GROUP (ORDER BY <i>sort_expression</i>)</code>	double precision[]	any sortable type	array of sort expression's type	multiple discrete percentile: returns an array of results matching the shape of the <i>fractions</i> parameter, with each non-null element replaced by the input value corresponding to that percentile

All the aggregates listed in **Table 9-51** ignore null values in their sorted input. For those that take a *fraction* parameter, the fraction value must be between 0 and 1; an error is thrown if not. However, a null fraction value simply produces a null result.

Each of the aggregates listed in **Table 9-52** is associated with a window function of the same name defined in **Section 9.21**. In each case, the aggregate result is the value that the associated window function would have returned for the "hypothetical" row constructed from *args*, if such a row had been added to the sorted group of rows computed from the *sorted_args*.

Table 9-52. Hypothetical-Set Aggregate Functions

Function	Direct Argument Type(s)	Aggregated Argument Type(s)	Return Type	Description
----------	-------------------------	-----------------------------	-------------	-------------

Function	Direct Argument Type(s)	Aggregated Argument Type(s)	Return Type	Description
<code>rank(<i>args</i>)</code> WITHIN GROUP (ORDER BY <i>sorted_args</i>)	VARIADIC "any"	VARIADIC "any"	bigint	rank of the hypothetical row, with gaps for duplicate rows
<code>dense_rank(<i>args</i>)</code> WITHIN GROUP (ORDER BY <i>sorted_args</i>)	VARIADIC "any"	VARIADIC "any"	bigint	rank of the hypothetical row, without gaps
<code>percent_rank(<i>args</i>)</code> WITHIN GROUP (ORDER BY <i>sorted_args</i>)	VARIADIC "any"	VARIADIC "any"	double precision	relative rank of the hypothetical row, ranging from 0 to 1
<code>cume_dist(<i>args</i>)</code> WITHIN GROUP (ORDER BY <i>sorted_args</i>)	VARIADIC "any"	VARIADIC "any"	double precision	relative rank of the hypothetical row, ranging from 1/ <i>n</i> to 1

For each of these hypothetical-set aggregates, the list of direct arguments given in *args* must match the number and types of the aggregated arguments given in *sorted_args*. Unlike most built-in aggregates, these aggregates are not strict, that is they do not drop input rows containing nulls. Null values sort according to the rule specified in the ORDER BY clause.

Table 9-53. Grouping Operations

Function	Return Type	Description
<code>GROUPING(<i>args</i>...)</code>	integer	Integer bit mask indicating which arguments are not being included in the current grouping set

Grouping operations are used in conjunction with grouping sets (see [Section 7.2.4](#)) to distinguish result rows. The arguments to the GROUPING operation are not actually evaluated, but they must match exactly expressions given in the GROUP BY clause of the associated query level. Bits are assigned with the rightmost argument being the least-significant bit; each bit is 0 if the corresponding expression is included in the grouping criteria of the grouping set generating the result row, and 1 if it is not. For example:

```
=> SELECT * FROM items_sold;
 make | model | sales
-----+-----+-----
  Foo  |  GT   |    10
  Foo  |  Tour  |    20
  Bar  |  City  |    15
  Bar  | Sport  |     5
(4 rows)
```

```
=> SELECT make, model, GROUPING(make,model), sum(sales) FROM items_sold GROUP BY ROLLUP(make,model);
 make | model | grouping | sum
-----+-----+-----+----
  Foo  |  GT   |         0 |   10
  Foo  |  Tour  |         0 |   20
  Bar  |  City  |         0 |   15
  Bar  | Sport  |         0 |    5
  Foo  |       |         1 |   30
  Bar  |       |         1 |   20
       |       |         3 |   50
(7 rows)
```



[Prev](#)

Range Functions and Operators

[Home](#)

[Up](#)

[Privacy Policy](#) | [Code of Conduct](#) | [About PostgreSQL](#) | [Contact](#)

Copyright © 1996-2022 The PostgreSQL Global Development Group

[Next](#)

Window Functions